



DAY 10 – Performance Optimization

INDIAN DATA CLUB

databricks

14 DAYS

AI CHALLENGE

DAY 10

Topic:
Performance Optimization

Challenge:

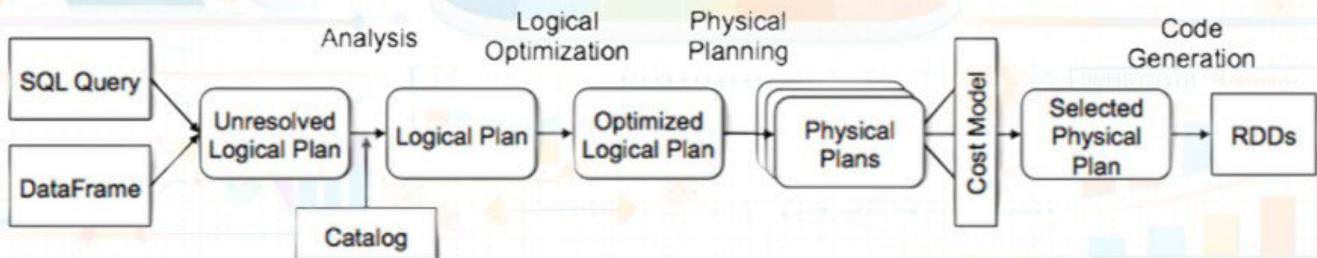
1. Analyze query plans
2. Partition large tables
3. Apply ZORDER
4. Benchmark improvements

#DatabricksWithIDC



What is a Query Execution Plan?

- A *query execution plan shows how Spark runs your query internally*
- *Includes:*
 - Table scans
 - Filters
 - Joins
 - Aggregations
 - Shuffles
- *Helps identify performance bottlenecks*





Why Query Plans Matter?

- ◆ Reveals why a query is slow
- ◆ Helps detect:
 - Full table scans
 - Expensive shuffles
 - Inefficient joins
- ◆ Essential for query tuning & optimization

How to Analyze Query Plans

SQL :

```
EXPLAIN FORMATTED SELECT * FROM  
table;
```

PySpark:

```
df.explain(True)
```

- Key things to look for:*
- Shuffle operations
 - Filter pushdown
 - Join strategies



What is Partitioning?

- ♦ *Partitioning splits a table into folders based on a column*
- ♦ *Example:*
 - `order_date=2024-01-01`
 - `order_date=2024-01-02`
- ♦ *Reduces data scanned during queries*

Benefits of Partitioning

- *Faster query performance*
- *Lower compute cost*
- *Efficient filtering on partition columns*
- *Improves scalability for large datasets*



Partitioning Best Practices

- ◆ *Use low to medium cardinality columns*
- ◆ *Best columns:*
 - Date, year, month
- ◆ *Avoid:*
 - High-cardinality columns (customer_id)
 - ◆ *Don't over-partition (small files problem)*



OPTIMIZE in Databricks

- *OPTIMIZE compacts many small files into fewer large files*
- *Improves read performance*
- *Reduces file management overhead*

```
OPTIMIZE sales_table;
```

What is ZORDER?

- *ZORDER clusters related data inside files*
- *Improves query performance for selective filters*
- *Works best on high-cardinality columns*

```
OPTIMIZE sales_table  
ZORDER BY (customer_id);
```



How to Cache Data

SQL:

```
CACHE TABLE sales_table;
```

PySpark:

```
df.cache()  
df.count()
```

Uncache when done:

```
UNCACHE TABLE sales_table;
```



When to Use Caching



Use caching for:

- Dashboards*
- Reused tables*
- Intermediate datasets*



Avoid caching:

- One-time queries*
- Very large tables*
- Memory-intensive workloads*